

Yang Xiao

918-906-9676 | yax3417@utulsa.edu | Google Scholar | GitHub | Personal Page

EDUCATION

University of Tulsa

Bachelor of Science in Computer Science

Tulsa, Oklahoma, USA

Aug. 2024 – May 2029 (Expected)

Focusing on Trustworthy and Efficient M.L. Model & System (Performance, Throughput, Latency), AI4Science.

Nankai University

Bachelor of Science in Applied Psychology

Tianjin, China

Sept. 2020 – June 2024

Learned Computer Science and Deep Learning from scratch by myself.

EXPERIENCE

Mayo Clinic

Data Science Intern

Rochester, Minnesota, USA

June 2025 – Aug. 2025

- Project 1: Fine-tuning LLM on clinical records and geographic data to predict the disease-related outcome of patients.
- Project 2: Fine-tuning LLM on drug information and gene sequential information to predict the sensitivity of a drug for the treatment.
- Both of them are under review of nature product journal AI.

University of Tulsa

Research Assistant

Tulsa, Oklahoma, USA

Aug. 2024 – Current

- We have published more than 10+ papers on the top-tier conferences and its relevant workshop range from ICLR, ICCV, CIKM, ECAI, and some workshops.

SELECTED PUBLICATION

- LiteFrame: Semantic-Aware Resampling for Efficient Micro-Video Recommendation *Under Review*
Yang Xiao, Huiyuan Chen, Kaiyuan Deng, Chao Jiang, Zinan Ling, Ruimeng Ye, Xiaolong Ma, Bo Hui
Impact: Developed a video preprocessing and training pipeline that reduces redundant computations by almost **99%**, significantly cutting down throughput & latency and improving recommendation performance.
- LightCache: Memory-Efficient, Training-Free Acceleration for Video Generation. *Under Review*
Yang Xiao, Gen Li, Kaiyuan Deng, Yushu Wu, Zheng Zhan, Yanzhi Wang, Xiaolong Ma, Bo Hui. [Paper]
Impact: Implemented a training-free acceleration framework for Diffusion models via 1) Asynchronous Cache Swapping. 2) Feature chunk. 3) Slicing latents to decode.
- Optimal Transport for Brain-Image Alignment: Unveiling Redundancy and Synergy in Neural Information Processing.
Yang Xiao, Wang Lu, Jie Ji, Ruimeng Ye, Gen Li, Xiaolong Ma, Bo Hui. *ICCV* 2025. [Paper]
Impact: Developed an Efficient Loss Function for multi-modality alignment via optimal transport.
- Efficient Knowledge Graph Unlearning with Zeroth-order Information
Yang Xiao, Ruimeng Ye, Bohan Liu, Xiaolong Ma, Bo Hui. *CIKM* 2025. [Paper]
Impact: Proposed an unlearning framework using ZO optimization. Avoided expensive Hessian-vector product computations, reducing GPU memory mostly from 85GB to **3GB** while maintaining the performance.
- Advancing Certified Robustness of Explanation via Gradient Quantization.
Yang Xiao, Zijie Zhang, Yuchen Fang, Da Yan, Yang Zhou, Wei-Shinn Ku, Bo Hui. *CIKM* 2024. [Paper]
Impact: Developed an explanation algorithm for Efficient Model Attribution via gradient quantization.

SERVICE

Journal: IEEE T-IFS. IEEE T-BD.

Conference: ECCV 2026, ICML 2026, CPAL 2026, LoG 2024

Workshop: NeurIPS workshop 2024, 2025

OTHER PUBLICATIONS & PRE-PRINTS

- **Yang Xiao**, et al. The Right to be Forgotten in Pruning: Unveil Machine... RegML@*NeurIPS* 2025 [Paper]
- **Yang Xiao**, et al. Knowledge Graph Unlearning with Schema. *LoG* 2024 *COLING* 2025 [Paper]
- **Yang Xiao**. Post-hoc and manifold explanations analysis of facial expression... *Undergrad Thesis* [Paper]
- R. Ye, Z. Wang, Z. Ling, **Yang Xiao**, et al. Your Language Model Secretly Contains... *ICLR* 2026 [Paper]
- R. Ye, Z. Wang, **Yang Xiao**, et al. Weak-to-Strong Generalization with Failure... *ICLR* 2026 [Paper]
- K. Deng, G. Li, **Yang Xiao**, et al. Forget Many, Forget Right: Scalable and Precise... *ICLR* 2026 [Paper]
- G. Li, **Yang Xiao**, et al. Sculpting Memory: Multi-Concept Forgetting in Diffusion... *ICCV* 2025 [Paper]
- B. Liu, **Yang Xiao**, et al. DBA-DFL: Distributed Backdoor Attacks with Network... *ECAI* 2025 [Paper]
- Z. Shi, **Yang Xiao**, et al. BMT-BENCH: A benchmark sports dataset for video... *ICIP* 2024 [Paper]
- R. Ye, Z. Wang, **Yang Xiao**, et al. Weak-to-Strong Generalization with... FoRLM@*NeurIPS* 2025. [Paper]
- Ruimeng Ye, **Yang Xiao**, et al. A Pilot Study of Weak-to-Strong... Bi-Align@*ICLR* 2025. [Paper]
- **Yang Xiao**, et al. LightCache: Memory-Efficient, Training-Free Acceleration for Video Generation. [Paper]
- M. Lee, **Yang Xiao**, et al. A Pharmacogenomic-Informed Representation Improves Multimodal... [Paper]
- B. Liu, Z. Zhang, P. He, Z. Wang, **Yang Xiao**, et al. A Survey of Lottery Ticket Hypothesis. [Paper]
- K. Deng, H. Zheng, M. Qing, K. Zhu, G. Li, **Yang Xiao**, et al. From Bits to Chips: An LLM-based... [Paper]